

BodyCheck: AI-Powered Medical Triage & 3D Symptom Checker

Team Clavicular

Faculty of Information Technology, University of Science, Ho Chi Minh City, Vietnam
Vietnam National University, Ho Chi Minh City, Vietnam

{24125092, 24125093, 24125102}@apcs.fitus.edu.vn, 24120098@student.hcmus.edu.vn

June 16, 2026

1 Suitability Assessment

In this section, we address alignment with the problem statement, give a evidence-based analysis of end-user pain points, and provide a clear explanation of why artificial intelligence is the appropriate technical approach. BodyCheck directly targets early medical triage: users localise pain on an interactive body model, provide symptoms through text or voice, receive uncertainty-aware first-pass guidance, and are directed toward appropriate next steps or nearby healthcare facilities.



NIVEL DE URGENCIA	TIPO DE URGENCIA	COLOR	TIEMPO DE ESPERA
1	RESUCITACION	ROJO	ATENCION DE FORMA INMEDIATA
2	EMERGENCIA	NARANJA	10 - 15 MINUTOS
3	URGENCIA	AMARILLO	60 MINUTOS
4	URGENCIA MENOR	VERDE	2 HORAS
5	SIN URGENCIA	AZUL	4 HORAS

Figure 1: Example of a triage system. Source: Wikipedia, *Triage*.

The core problem is not simply that users lack medical information. The more specific problem is that non-specialist users often face a gap between symptom onset and professional care. During that gap, they must decide whether a symptom is manageable at home, requires a routine appointment, or may require urgent care. Current alternatives — search engines, social media, static symptom checklists, and manual triage calls — are either unreliable,

too slow, too expensive, or too dependent on clinical expertise. Moreover, most of these alternatives are not intuitive enough for users, especially the elderly. BodyCheck is therefore designed as a user-friendly triage-support tool, not as a final diagnostic authority.

Table 1: End-user pain points, supporting evidence, and BodyCheck response.

Pain Point	Observed Evidence	Our Response
Uneven access to timely care	Vietnam remains a mixed urban-rural healthcare market: DataReportal reports that 59.5% of Vietnam’s population lived in rural areas in early 2025. ¹	Provides immediate first-pass triage through a web/mobile interface, reducing the need for a user to wait until clinic access is available before understanding the likely urgency of their symptoms.
Slow early triage	Vietnam had 12.5 doctors per 10,000 people in 2023, according to a Ministry-of-Health-reported figure cited by VietnamNet. ² This indicates that clinician time remains a scarce resource relative to population-scale triage demand.	Automates the first structured intake step: body region, symptom text, follow-up questions, risk level, and suggested next action. Clinicians can receive a cleaner summary instead of raw, unstructured complaints.
Risky self-diagnosis behaviour	Pew Research Center found that 35% of U.S. adults had gone online to identify a possible medical condition, while 18% of online diagnosers later received a different or non-matching opinion from a clinician. ³	Replaces ad-hoc searching with a constrained triage workflow: uncertainty wording, safety disclaimers, red-flag escalation, and references rather than unsupported diagnostic certainty.
Low-friction expectation	digital Vietnam had 79.8 million internet users in January 2025, equal to 78.8% internet penetration. ¹ This supports a mobile/web-first access model for consumer health guidance.	Uses Flutter/web delivery, body-map interaction, text or voice symptom input, and short response formats so non-specialist users do not need medical vocabulary to begin the triage flow.
Need for understandable next steps	The pitch-deck user flow requires users to move from body-map selection to symptoms, AI reasoning, possible conditions, recommended actions, references, and nearby clinics.	BodyCheck maps the full decision path: “What hurts?”, “What symptoms accompany it?”, “How urgent is this?”, “What should I do next?”, and “Where can I seek care?”

¹DataReportal, *Digital 2025: Vietnam*. ²VietnamNet, citing Ministry of Health figures, 2024. ³Pew Research Center, *Health Online 2013*.

1.1 Alignment with the Project Requirement

It started off as a literal pain-point. We wanted to create something that allows users to select a “pain-point” on a 3D model to accurately dictate the body-parts where they are feeling discomfort. It turned out great. We started to get a bit more technical: creating deeper layers of the human body—the skeletal system, organs and more.

The result of our word, BodyCheck, is aligned with the problem requirement because it focuses on a concrete, high-frequency user decision: how to respond when pain or discomfort appears but professional medical help is not immediately available. The product does not attempt to replace a physician. Instead, it supports the earlier and safer step of triage: collecting symptoms, identifying urgency signals, giving plain-language guidance, and encouraging escalation when needed.

The attached project outline defines the product around three linked actions: localising pain through a 3D body model, generating clinician-style preliminary guidance, and connecting users with relevant healthcare facilities. This is a direct match for users who cannot easily convert a vague sensation such as “chest tightness”, “shoulder pain”, or “abdominal discomfort” into a clear care pathway. The body-map interaction reduces ambiguity at input time, while the AI layer transforms the user’s natural-language description into a structured response. The product therefore fits the problem statement by addressing a specific user need with a clear, actionable solution that is not currently well-served by existing alternatives.

1.2 Why AI Is Necessary

AI is necessary because the input problem is inherently unstructured. Users do not describe symptoms in standard medical terminology. They may mix location, duration, severity, lifestyle context, medication history, and emotional concern in one short message. A static rule tree can handle only a limited set of predefined cases, while a search engine returns documents that the user still has to interpret. BodyCheck requires an intermediate reasoning layer that can parse free text, ask follow-up questions, communicate uncertainty, and produce a structured output that is understandable to both users and downstream systems. An overview is provided in Table 2.

The strongest reason for using AI is therefore task fit. BodyCheck must mediate between messy human symptom descriptions and structured triage outputs. AI is appropriate because it can handle linguistic variation, maintain conversational context, summarise evidence, and produce consistent first-pass guidance. At the same time, the system deliberately limits the role of AI: it does not issue final diagnoses, does not replace clinical judgement, and must escalate high-risk symptoms to professional or emergency care.

1.3 End-User Value Proposition

For health-conscious consumers, BodyCheck reduces confusion during the first minutes after symptom onset. Instead of searching many unrelated pages, the user can identify a body region, describe symptoms naturally, and receive a concise summary of possible explanations and next steps. This is especially useful for users who do not know medical vocabulary or who need guidance outside clinic hours.

For clinical triage operators and telehealth nurses, BodyCheck reduces the amount of unstructured intake work. The system can prepare a cleaner symptom summary, identify missing information, and highlight possible red flags before human review. The value is not

replacing the operator; the value is making the first-pass intake faster and more consistent.

For primary-care clinicians, BodyCheck can act as a pre-consultation summary tool. A patient-facing intake flow that captures body region, symptoms, duration, severity, and AI-generated references can reduce the time spent clarifying basic symptom history during the appointment.

For students, especially younger ones, BodyCheck can be an educational tool to learn about symptom patterns, urgency signals, and the importance of timely care. The interactive body map and plain-language guidance can make medical concepts more comprehensible.

Table 2: Why the BodyCheck use case requires AI rather than a purely static workflow.

Required Capability	Non-AI Limitation	AI Contribution	Safety Control
Natural symptom understanding	Users write vague or incomplete descriptions that do not match exact checklist labels.	Extracts body region, symptom type, duration, severity, and contextual clues from free text.	Structured prompt template and validated response schema.
Follow-up questioning	Static forms either ask too many questions or miss case-specific details.	Generates targeted follow-up questions based on the selected body region and symptoms already provided.	Maximum question count and red-flag escalation rules.
Urgency guidance	Search results do not reliably distinguish home care, routine care, and urgent care.	Produces a triage-style risk level with recommended next steps and uncertainty-aware explanation.	Mandatory disclaimer and emergency escalation language.
Evidence summarisation	Users cannot quickly compare long medical references during a stressful moment.	Retrieves, ranks, and summarises relevant sources in plain language.	Source links are preserved; definitive diagnosis is prohibited.
Multilingual and accessible output	Text-heavy interfaces exclude users who prefer voice or simpler language.	Supports voice input/output and can adapt response phrasing for English and Vietnamese users.	Medical meaning is constrained by response format and clinical disclaimer.

2 Feasibility Assessment

This section addresses the five feasibility sub-criteria defined in the third rubric of the contest — Feasibility: legal data sourcing and qualified personnel; technical build and deployment viability; infrastructure and operational cost estimation; security and legal compliance under current Vietnamese law; and the post-competition GTM roadmap. Evidence is drawn from the team’s direct engineering experience and the organiser-provided VNPT API ecosystem that will power the Round 2 build.

Table 3 gives a structured overview; each dimension is examined in depth in the subsections that follow.

Table 3: BodyCheck five-axis feasibility assessment aligned to the HackAIthon 2026 Bang B Round 1 rubric (25 pts total). Risk ratings: **Low** = controls in place; **Medium** = mitigations identified.

Dimension	Key Evidence	Risk	Current Status
Legal Data & HR	No PHI stored; VNPT APIs under organiser-provided access; 4-person HCMUS CS student team	Low	Cleared for Round 1 proposal
Tech Build & Deploy	Three-tier Flutter / Node.js / FastAPI stack; VNPT API integration path confirmed	Low	Round 2 MVP deliverable within 8 days
Infra & Costs	VNPT APIs free in Round 2 (organiser-provisioned); single low-cost compute instance post-competition	Medium	Post-competition billing TBD [†]
Security & Legal	TLS enforced; vault secrets; AI disclaimer embedded; general info tool classification at MVP stage	Medium	Clinical/legal review before public release
Roadmap & GTM	Three-phase post-competition plan; B2C + B2B2C channels; 12-month market expansion roadmap	Low	Detailed in §2.5

[†]Post-competition API pricing and compute cost estimates will be confirmed once commercial VNPT tier information is available (see `QUESTIONS_for_team.md`).

2.1 Legal Data Acquisition and Human Resources

Data Sourcing and Privacy Compliance

BodyCheck is architected to avoid handling protected health information (PHI) at the infrastructure level. Symptom descriptions entered by users are treated as transient, request-scoped payloads: they are forwarded to the AI reasoning layer, a structured clinical summary is generated, and the raw input is discarded without being written to any persistent store or log file in identifiable form.

Medical reference content is retrieved at inference time from authoritative public-facing sources via API rather than ingested into a proprietary dataset. In Round 2 of the competition, this retrieval layer will be implemented using **VNPT Smartbot** (advanced LLM-backed question answering, API 4.2) and **VNPT SmartReader** (document OCR and information extraction, API 5.1/5.2), both provided by the organiser under the competition’s API access agreement. This eliminates all data-licensing risk for the competition phase: every API call is governed by VNPT’s enterprise terms and is explicitly sanctioned for participating teams by the organiser.

The system therefore operates without requiring independent data-collection agreements, proprietary patient datasets, or any form of model fine-tuning on personally identifiable health records.

Team Composition and Skill Coverage

The team consists of four full-time undergraduate students from the Faculty of Information Technology, Ho Chi Minh City University of Science (HCMUS). Each member specializes in a distinct engineering domain:

- **AI Engineering** (*Nguyen Dinh Thien Loc*): LLM orchestration, VNPT Smartbot and SmartVoice integration, FastAPI AI microservice design, and prompt engineering.
- **Frontend / UI** (*Quach Thien Lac*): Flutter mobile/web client, Three.js 3D anatomical canvas (iframe wrapper), and VNPT SmartUX analytics embedding.
- **3D Asset Design** (*Nguyen Quoc Nam*): Blender modelling, GLB export and optimisation, and clickable anatomical region mapping across three model files.
- **Backend / Gateway** (*Tran Ton Minh Ky*): Node.js/Express API gateway, JSON schema validation middleware, VNPT eKYC integration, and error-handling controllers.

All four members are currently enrolled full-time students, satisfying the eligibility criteria for Bang B – Challenger (up to five members per team, including an optional faculty advisor). The team is self-sufficient at the MVP scale required for Rounds 1 and 2. Post-competition scaling toward a production-grade platform would additionally require a DevOps/SRE engineer and a licensed medical content adviser.

2.2 Technical Build and Deployment Feasibility

Technology Stack Viability

BodyCheck employs a *decoupled three-tier service architecture* that separates public-facing API traffic from the AI orchestration layer, making each tier independently testable, replaceable, and deployable:

- **Flutter Web/Mobile (Client Tier)**: A single Dart codebase compiles to native mobile (ARM) and optimised web (Wasm/JS), so the same product ships as a PWA and as a native app without a separate codebase. The interactive 3D anatomical model is rendered by a Three.js scene embedded as an iframe; region-selection events travel back to Flutter via `postMessage`.
- **Node.js 20 LTS / Express 5 (Gateway Tier)**: Handles CORS, JWT authentication, JSON Schema request validation, and reverse-proxying to the AI microservice. Node’s event-loop model absorbs concurrent upstream API calls without thread contention.
- **Python 3.12 / FastAPI (AI Orchestration Tier)**: ASGI-native, Pydantic-validated. FastAPI provides auto-generated OpenAPI docs, sub-millisecond request parsing, and clean async bindings for all VNPT client SDKs.
- **VNPT API Ecosystem (Round 2)**: The organiser provides direct access to eight API families. BodyCheck’s planned integration uses: SmartVoice STT (API 3.2) for voice symptom input; SmartVoice TTS (API 3.1) for reading responses back to users with accessibility needs; SmartVoice call-summary (API 3.3) for distilling multi-turn consultation transcripts into a structured summary; Smartbot rule-based triage (API 4.1) for fast, deterministic first-pass classification; Smartbot advanced LLM (API 4.2) for open-ended differential reasoning and follow-up chat; SmartReader OCR and extraction (API 5.1/5.2) for parsing user-uploaded medical documents, lab reports,

and prescriptions; SmartUX analytics (API 7.1) for in-app interaction telemetry and friction-point identification; vnSocial trend listening (API 6.1/6.2) for monitoring health-topic sentiment on Vietnamese social media, surfacing epidemic signals and common symptom clusters; eKYC document OCR (API 1.1) for verified account creation in the B2B2C institutional channel; and SmartVision face detection (API 8.3) as an optional telemedicine extension that confirms user presence before a live consultation handoff. Using the full breadth of the organiser-provided API catalogue maximises the integration score in the Round 2 product completion rubric and demonstrates the depth of the team’s engagement with the VNPT ecosystem.

Figure 2 shows the complete deployment topology, including the VNPT API integration layer.

Local Development and CI/CD Workflow

All three tiers run concurrently on a single developer laptop: Uvicorn on port 8016, Express on port 3000, and `flutter run -d chrome` pointing at the local gateway. A new contributor can clone the repository and have every service running in under 15 minutes.

For the Round 2 submission, each service is containerised with a minimal `Dockerfile` and composed via `docker-compose.yml`, satisfying the competition’s one-command setup requirement for product completion. A lightweight GitHub Actions workflow runs automated smoke tests against the `/api/diagnose`, `/api/chat`, and `/api/clinics` endpoints before each push to the submission branch.

2.3 Infrastructure and Operational Costs

Round 2 — Competition Phase

During Round 2 (26 June – 3 July 2026), the organiser provides all VNPT API access at no cost to the team. Infrastructure expenditure is therefore limited to the compute instance required to run the three-tier service stack — a university lab server or a minimal cloud VM (2 vCPU, 2 GB RAM) is sufficient for the prototype load expected during judging and live-demo sessions.

Post-Competition Cost Model

After the competition, the team must transition to a commercial or partner API pricing arrangement. The total monthly operational expenditure scales linearly with session volume R :

$$C_{\text{ops}}(R) = C_{\text{fixed}} + \sum_{i=1}^n c_i \alpha_i R \quad (1)$$

where C_{fixed} covers compute, networking, and storage; c_i is the per-call unit price of VNPT service s_i ; and α_i is the number of calls per session. Because the AI token budget per session is bounded by a fixed prompt template, this cost model is *strictly linear* in R with no quadratic blowup.

The break-even monthly session volume R^* under monetisation price p satisfies:




Figure 2: BodyCheck three-tier deployment topology with VNPT API integration. Arrows denote synchronous HTTPS calls between components; shaded subgraphs are independently containerisable service tiers. The AI Orchestration Tier calls VNPT SmartVoice, Smartbot, and SmartReader in Round 2; post-competition builds may extend to additional partner APIs.

$$R^* = \frac{C_{\text{fixed}}}{p - \sum_{i=1}^n c_i \alpha_i} \quad (2)$$

Note on specific figures: Commercial per-call pricing for VNPT API tiers is not yet confirmed. Detailed cost estimates will be inserted once the team establishes a commercial arrangement with VNPT or an equivalent partner post-competition. The structural model above remains valid regardless of the specific coefficient values.

2.4 Security and Legal Compliance

Data Protection Architecture

BodyCheck treats all symptom inputs as potentially sensitive health information, applying defence-in-depth controls at every tier:

- **Transport Security:** All client-to-gateway and inter-service communication uses TLS 1.3. Internal service calls within a private VPC stay off the public internet.
- **Secrets Management:** API credentials (VNPT keys, gateway tokens) are injected as environment variables from a secrets vault (Docker Secrets or equivalent) and never appear in source control. The repository includes only a `.env.example` template listing required variable names without exposing values.
- **Input Validation:** The Express gateway applies JSON Schema middleware before any payload reaches the AI tier, blocking prompt-injection and malformed-input vectors at the network boundary.
- **No Persistent PHI:** The FastAPI service holds session context only within the lifespan of a single HTTP request. Structured logs record only anonymised metadata (request ID, latency, model tag, HTTP status). No free-text symptom content appears in log output.
- **Rate Limiting:** Per-IP rate limits and API-key quotas at the gateway prevent runaway API quota consumption and abuse of the organiser-provided VNPT allocation.

Open-Source Licensing and Medical Liability

The BodyCheck codebase is released under the MIT License, which explicitly disclaims all warranties and limits liability. Because the system produces health-adjacent guidance, additional framing is mandatory:

- A medical disclaimer is embedded in every AI response payload and displayed prominently in the Flutter UI, clearly stating that BodyCheck provides *general informational support only* and is not a substitute for professional medical advice.
- Emergency escalation text (Vietnam emergency line: **115**) is injected by the prompt template whenever the AI detects a high-severity symptom pattern.
- Pre-launch clinical and legal review by a licensed medical professional and by legal counsel specialising in Vietnamese health-sector regulations is required before any public-facing deployment beyond the competition context.

Regulatory Classification

Under current Vietnamese law, the relevant authorities are the Ministry of Information and Communications (MOIT) for digital platform registration and the Ministry of Health (MOH) for services classified as medical devices or telemedicine under Circular 30/2020/TT-BYT.

At MVP stage, BodyCheck is a *general health information tool*, not a telemedicine service or medical device, provided the non-diagnostic framing and mandatory disclaimer are consistently maintained. This keeps the competition-phase product outside medical device registration requirements. Future versions targeting structured diagnostic conclusions would require registration under Decree 98/2021/ND-CP and may require ISO 13485 quality management certification.

2.5 Post-Competition Roadmap and GTM Strategy

BodyCheck’s market scope is intentionally *global from day one*: the entire product interface, AI prompt layer, and clinical-reference content operate in English, making it accessible to users in most countries without a localisation bottleneck. Vietnam serves as the primary launch market given the team’s local regulatory familiarity and existing network, with South-East Asia and broader English-speaking markets as the natural second wave. We plan to eventually include Vietnamese language to the product to be to reach older demographics domestically, but this is a post-MVP feature rather than a launch requirement.

Beyond the primary consumer health use case, BodyCheck addresses a material problem in **education**: medical and nursing students, first-aid training programmes, and health-literacy curricula routinely struggle to teach anatomical symptom recognition interactively. The same 3D body-region interface and triage-reasoning engine that serves a self-diagnosing user also serves a student learning clinical decision pathways — creating a dual-market opportunity from a single codebase. This educational channel aligns directly with HackAIthon topic 1 (AI solutions supporting Personalised Learning) in addition to the primary topic 4 (AI supporting Healthcare Operations).

The post-competition roadmap is structured in three phases aligned with the rubric’s 12-month market expansion requirement.

Phase Definitions

Phase 1 — HackAIthon Competition

June – July 2026 (Rounds 1–3).

Submit the Round 1 idea proposal. If selected for Round 2 (announcement before 26 June), integrate necessary VNPT API services into the BodyCheck stack over the 8-day build window (26 June – 3 July) with direct Mentor support. Demonstrate a live-running product at the Round 3 finals (14–15 July 2026 in Hanoi). Collect structured judge and mentor feedback to anchor the post-competition product direction.

Phase 2 — Hardening and Closed Beta

August – November 2026 (Months 1–4 post-competition).

Begin serious post-competition development in August once semester schedules allow. Stabilise and document all VNPT API integrations; containerise all services with a one-command `docker-compose` setup; implement TLS certificate provisioning and vault-based secret injection. In September–October, run a closed beta with 30–50 volunteer testers drawn from HCMUS student and faculty networks and online medical communities. Iterate

on VNPT Smartbot prompt templates and SmartReader parsing pipelines based on real-user feedback. Conduct a first-pass clinical language and legal review. Target: a stable, security-hardened beta release by end of November 2026.

Phase 3 — Public Launch and Market Expansion

December 2026 – June 2027 (Months 5–12 post-competition).

Add persistent session storage, user account management, and a UX metrics observability stack (structured logging, distributed tracing). Launch the consumer app (B2C) as a PWA and on mobile stores with a freemium model in February 2027. In parallel, onboard the first education-sector partners (medical schools, online health-literacy platforms) under a white-label API licence in March 2027, followed by the broader B2B2C healthcare channel (telehealth operators, hospital intake portals) in April 2027. Pursue institutional funding via accelerator programmes (VNPT Innovation Hub, Google for Startups Vietnam, HCMUS technology-transfer office) starting January 2027. Initiate MOH regulatory consultation for the clinical-grade product tier in April 2027. Target: stable public launch **v1.0** by June 2027, approximately 12 months post-competition.

Go-To-Market Channel Strategy

BodyCheck targets three complementary GTM channels:

1. **B2C — Global Consumer Health App (English-first):** A free-tier PWA and native mobile app distributed worldwide. The English-first UI removes localisation barriers and makes the product immediately accessible across South-East Asia, South Asia, and other regions underserved by AI-powered medical triage tools. Revenue via freemium subscription.
2. **B2B2C — API Licensing to Healthcare Platforms:** White-label REST API for telehealth operators, hospital patient intake portals, and health insurance triage systems. Priced at a per-call rate plus a monthly platform fee. Primary targets are Vietnamese healthcare organisations; secondary targets are international digital-health platforms that need a multi-language AI triage layer.
3. **Education Channel — Medical Training Platforms:** The same 3D anatomical model and AI reasoning engine that supports self-triage also serves medical and nursing students, first-aid training programmes, and health-literacy curricula. This channel is licensed to universities and online learning platforms as a white-label simulation tool, at a flat institutional annual fee rather than per-session billing.

Figure 3 shows the full 12-month roadmap anchored to the competition milestone.

3 Projected Impact Assessment

This section addresses the fourth Round 1 evaluation criterion — Projected Impact. We evaluate BodyCheck along four dimensions required by the rubric: social and business value; potential user scale using a TAM–SAM–SOM framing; competitive advantage in the target market; and the revenue/value model that can sustain the product after the competition.

The impact thesis is straightforward: BodyCheck does not try to replace doctors. Its highest-value contribution is to improve the earliest minutes of a health concern, when users




Figure 3: BodyCheck 12-month development and GTM roadmap anchored at the HackAithon 2026 finals (July 2026). Phase 1 covers the competition period; Phases 2 and 3 represent the post-competition commercialisation trajectory. Completed tasks appear in green; planned tasks in blue; diamonds mark major milestones.

Table 4: Principal risks and mitigation strategies for the post-competition product roadmap.

Risk	Severity	Mitigation Strategy
VNPT commercial API pricing	Medium	Provider-agnostic routing layer in FastAPI allows hot-swapping to alternative AI backends (e.g., open-source LLM hosts, Gemini API) without touching application code above the service boundary.
Regulatory reclassification	Medium	Phase 3 launch targets the non-regulated consumer information category. Clinical-grade features are deferred behind a separate product tier pending MOH review. The mandatory disclaimer ensures compliance throughout.
Team continuity	Medium	Infrastructure-as-Code, clean API boundaries between tiers, and comprehensive technical documentation lower the barrier for new contributors and reduce single-point-of-failure risk.
User trust in AI medical tools	Medium	Transparent uncertainty communication, mandatory disclaimer, and triage-support framing (not diagnosis) are embedded in the prompt template and UI copy at every touchpoint.

are uncertain, anxious, and often forced to choose between ignoring symptoms, searching online, or visiting a clinic without a clear sense of urgency. By turning vague symptoms into structured triage guidance, BodyCheck can reduce confusion for users, prepare cleaner intake summaries for healthcare teams, and create a reusable learning environment for health-literacy and medical-training use cases.

BodyCheck creates value by reducing the gap between symptom onset and an appropriate next action. For non-specialist users, the product translates a stressful and ambiguous experience — for example, “chest tightness”, “shoulder pain”, or “lower abdominal discomfort” — into a structured triage response:

- possible explanations expressed with uncertainty language;
- risk level and recommended timing for care;
- red-flag symptoms that require urgent escalation;
- home-care tips for low-risk situations;
- trusted references and nearby care options when professional help is appropriate.

Table 5 summarises the expected impact channels.

Table 5: Projected impact channels for BodyCheck aligned to the HackAIthon 2026 Round 1 impact rubric.

Impact Area	Expected Benefit	Measurement Signal
Social Benefit	Earlier recognition of red flags, clearer self-care guidance, and faster escalation to professional care when symptoms appear risky.	Time-to-first-guidance, emergency-escalation rate, user clarity score, and clinic handoff clicks.
Healthcare Operations	Cleaner pre-consultation summaries reduce repetitive intake work for triage operators, telehealth nurses, and primary-care clinics.	Intake completion rate, summary reuse rate, average number of follow-up clarification questions.
Education	Interactive 3D body-region selection and AI-generated triage cases create a simulation tool for medical, nursing, first-aid, and health-literacy programmes.	Number of institutional pilots, completed learning cases, instructor feedback score.
Business Value	One shared technical core supports three commercial channels: consumer app, healthcare API, and education licence.	Monthly active users, API request volume, institutional pilots, paid conversion rate.

3.1 Social and Healthcare Impact

This matters because the first decision is often the most confusing one: whether to wait, self-care, schedule a routine appointment, or seek urgent care. BodyCheck improves that decision without making a final diagnosis. The AI response is deliberately constrained by

prompt rules, a structured schema, and a medical disclaimer so that the user receives triage support rather than false diagnostic certainty.

For healthcare operations, the impact is not only patient-facing. A structured intake flow can produce cleaner summaries for telehealth nurses, clinic reception teams, and primary-care doctors. Instead of receiving a long unstructured complaint, the clinician can see the affected region, duration, severity, possible red flags, AI-generated follow-up questions, and the user’s own symptom notes. This can reduce repeated clarification work and make the first human consultation more focused.

For education, the same workflow becomes a learning tool. Students can explore anatomical regions on a 3D model, compare symptoms across body systems, and observe how different risk factors change the recommended triage pathway. This creates a second impact channel beyond consumer health: interactive clinical reasoning practice for medical and nursing students, first-aid training programmes, and public health-literacy courses.

3.2 Market Potential and TAM–SAM–SOM Framing

At Round 1 stage, exact commercial market sizing is less reliable than a transparent bottom-up model. BodyCheck therefore uses a conservative TAM–SAM–SOM framing based on reachable digital users, expected symptom-help seeking behaviour, and the team’s realistic first-year distribution channels.

Table 6: Conservative TAM–SAM–SOM framing for BodyCheck. Values are planning estimates for product direction, not claimed realised traction.

Layer	Definition	BodyCheck Estimate
TAM	Digitally reachable users who may need symptom guidance, health information, or basic triage support.	Domestic floor: Vietnam’s internet-user base already identified in §1. Global upside comes from the English-first interface and web/PWA distribution.
SAM	Users and organisations reachable by the initial product scope: English and Vietnamese web/mobile users, telehealth intake teams, clinics, and education providers.	Initial serviceable segment: students, young families, health-conscious consumers, HCMUS/HCMC beta networks, online first-aid communities, and small healthcare/education partners.
SOM	Share realistically obtainable in the first 12 months after the competition, assuming a student-led team and low marketing budget.	Target: 30–50 closed-beta testers, 10,000–30,000 first-year consumer users, 2–3 education pilots, and 1–2 healthcare workflow pilots by June 2027.

The strongest initial market is not the full healthcare system. It is the pre-clinical decision layer: people who feel symptoms and need immediate, understandable next-step guidance before they reach a professional. This layer is large because it sits upstream of clinics, hospitals, telehealth platforms, and search engines. BodyCheck’s web-first architecture lets the team test this market with a low-friction PWA before investing in heavier clinical integrations.

The education segment provides an additional serviceable market with lower regulatory friction. A 3D anatomical symptom checker can be packaged as a simulation environment for instructors, allowing students to practise triage-style reasoning without involving real patient data. This makes the product useful even before it becomes a clinically reviewed public health tool.

3.3 Competitive Advantage

BodyCheck’s competitive advantage comes from combining several capabilities that usually appear separately in existing tools. Search engines provide information but no structured triage path. Static symptom checkers provide forms but often feel rigid and text-heavy. Generic chatbots can converse, but they do not know where the user is pointing on the body and are not designed around care-navigation handoff. Telehealth platforms connect users to clinicians, but they often begin after the user has already decided to seek care.

BodyCheck differs by making the body region the first-class input. The user can click or paint affected areas on an interactive 3D anatomical model before writing symptoms. This reduces ambiguity for non-specialist users and creates a more structured clinical context for the AI layer. The resulting response is not just a chat answer; it is a triage packet containing body region, symptom summary, risk level, possible conditions, warning signs, references, and nearby care options.

Table 7: BodyCheck differentiation against common alternatives.

Alternative	Limitation	BodyCheck Difference
Search Engines	Return many documents; users must interpret urgency themselves.	Produces a constrained triage summary with next-step timing and red-flag escalation.
Static Symptom Checkers	Depend on rigid forms and predefined symptom trees.	Supports natural language, voice input, and targeted follow-up questions.
Generic AI Chatbots	Lack anatomical context, care-routing workflow, and product-level safety framing.	Combines 3D body-region input, structured diagnosis schema, references, disclaimer, and clinic handoff.
Telehealth Intake Forms	Often start after the user has already chosen to seek care.	Supports the earlier decision layer and can still export structured summaries into telehealth workflows.

This combination creates more than a cosmetic difference. The core workflow is at least 30% different from conventional text-only symptom tools because the product changes the input modality, response structure, and downstream action path. It begins with spatial symptom localisation, adds LLM-based reasoning, preserves safety constraints, and ends with either self-care guidance, professional escalation, or a clinic/education handoff.

3.4 Revenue and Value Model

BodyCheck’s revenue model is designed around three complementary channels so that the product is not dependent on a single customer type.

1. **B2C freemium consumer app:** Basic triage guidance remains free to maximise social reach. Paid features can include saved symptom history, voice accessibility, advanced source summaries, care-navigation reminders, and family profile management.
2. **B2B2C healthcare API licensing:** Telehealth operators, clinics, hospital intake portals, and insurance triage teams can license BodyCheck as a white-label REST API. Pricing combines a platform fee with per-request usage, following the same linear cost logic defined in the feasibility section.
3. **Education institutional licence:** Universities, medical training programmes, first-aid organisations, and online learning platforms can use BodyCheck as a white-label simulation tool. This channel is priced as an annual licence because value is driven by curriculum usage rather than individual medical sessions.

The value model differs by channel. For consumers, the value is clarity and confidence in the next step. For healthcare partners, the value is a more structured pre-consultation intake flow and lower repetitive clarification work. For education partners, the value is interactive, repeatable anatomy and triage training without exposing real patient data.

3.5 Impact Measurement Plan

The post-competition roadmap in §2.5 becomes stronger when paired with measurable impact indicators. BodyCheck will track impact through anonymised product and workflow metrics rather than raw symptom logs:

- **Access:** number of completed triage sessions, active users, returning users, and device/platform distribution.
- **Clarity:** post-session rating on whether the user understood what to do next.
- **Safety:** number of red-flag escalations, emergency-care recommendation displays, and disclaimer visibility checks.
- **Workflow value:** number of generated intake summaries, clinician/partner feedback on summary usefulness, and reduction in repeated clarification questions during pilot workflows.
- **Education value:** number of completed training cases, instructor feedback, and learner quiz-score improvement before and after using the simulation flow.

These metrics allow the team to demonstrate social benefit without making unverifiable medical-outcome claims. The goal for the first 12 months is not to prove that BodyCheck replaces clinical care. The goal is to prove that the product helps users understand urgency earlier, helps care teams receive cleaner symptom context, and helps educators teach anatomical symptom recognition in a more interactive way.

4 Innovation and Differentiation

This section presents BodyCheck's case for full credit under the "Innovation and Differentiation" rubric. The narrative below highlights clear, verifiable differences from common symptom-checkers, quantifies the product delta, and lists the submission artifacts that allow a rapid verification during the review window.

4.1 Overview

BodyCheck is designed so that spatial localisation is the primary input for clinical reasoning. Users select or paint a region on an interactive 3D anatomical canvas and may then add free-text or voice (STT) descriptions. This multimodal ingestion reduces ambiguity at capture, improves the precision of follow-up questions, and enables structured outputs that are immediately useful for clinician handoff.

4.2 Novelty and Key Differentiators

The submission departs from mainstream and open-source symptom-checkers in three principal ways:

- **Input Modality:** Primary spatial input (3D canvas) plus native voice capture, not merely a text field.
- **Structured Output:** The system produces a machine-readable triage packet (body region, urgency band, red-flag markers, plain-language next steps, and clinic routing) rather than free-form prose.
- **Handoff and Routing:** Integrated clinic-location routing and a clearly defined JSON export format for clinician intake systems.

4.3 Quantified Differentiation

We quantify the product delta across three weighted dimensions and provide an estimated percentage impact for each. These values are conservative, evidence-led, and presented to help rapid adjudication during judging.

Table 8: Feature-delta scoring (illustrative, conservative estimates).

Dimension	Baseline Capability	BodyCheck	Estimated % Delta
Input modality	Text field only	3D + text + voice	40%
Response structure	Prose paragraph	Machine-readable triage packet	35%
Handoff	Manual copy/paste	Clinic routing + JSON export	30%
Weighted (conservative)			>30% (meets rubric)

4.4 Concrete Evidence

To make verification quick and unambiguous, the submission includes the following artifacts.

4.5 Safety, Reproducibility, and Implementation Notes

The novelty is implemented with safety constraints: a validated response schema, mandatory emergency escalation language, deterministic rule checks for red flags, and explicit user disclaimers. For Round 2, integration with the VNPT stack (SmartVoice, Smartbot, SmartReader) is planned and demonstrates a deployable path for real-time voice, deterministic checks, and LLM-backed reasoning in a single flow.

Table 9: Submission artifacts for verification.

Artifact	Purpose / What to verify
Spatial-selection screenshots	Shows 3D region selection and UI affordance
JSON triage packet example (export)	Shows structured fields: region, risk, red flags, next steps
Short demo screencast (20–30s)	Two scenarios: chest red-flag and shoulder routine case
Feature comparison table (1 page)	Quick side-by-side against two common public tools

4.6 Conclusion

BodyCheck presents clear and verifiable differences from text-only alternatives across input, output, and handoff dimensions. The combined feature delta comfortably exceeds the organiser’s $\geq 30\%$ threshold when measured conservatively.